

Load Balancing across Microservices

Yipei Niu¹, Fangming Liu¹, Zongpeng Li²



Email: fmliu@hust.edu.cn

¹Cloud Datacenter & Green Computing/Communications Research Group


¹Huazhong University of Science & Technology

²University of Calgary

Let's hail a cab




Toward monolithic hell



- The **monolithic** architecture
 - Services and modules are tightly coupled


Complex to maintain, debug, and evolve!

Breaking monolithic software into microservices



A development technique that structures an **monolithic** application as a collection of **loosely** coupled services


Breaking monolithic software into microservices



- Microservice has been adopted by Amazon, Netflix, and Uber


Source: <https://thenewstack.io/led-amazon-microservices-architecture/>
<https://www.nginx.com/blog/microservices-at-netflix-architectural-best-practices/>
<https://eng.uber.com/soa/>

Breaking monolithic software into microservices



- Each microservice has **limited** functionalities
- Instances runs **independently** on containers
- Requests are served by microservice chains


Breaking monolithic software into microservices



- Each microservice has **limited** functionalities
- Instances runs **independently** on containers
- Requests are served by microservice chains

Imbalance load across microservices


- Workload is fluctuating
- Different QoS of chains
- Different service time in microservices



We need to balance load from perspective of chains!

Imbalance load across microservices

- Workload is fluctuating
- Different QoS of chains
- Different service time in microservices



How to identify chains?

How to determine instance scale of a chain?

Tackling the complexity

- Typical load balancer
 - Haproxy, Nginx

Fail to identify different chains

- Communication pattern
 - Application-layer protocols
 - HTTP for RESTful API
 - AMQP* for message queue

Complicated to operate


- Large scale
 - Netflix employs over 600 microservices

An efficient strategy is required

*AMQP: Advanced Message Queueing Protocol

Abandon load balancer


- Replace load balancer with message queue
- Upstream instances declare a channel
- Downstream instances are aware of what messages (requests) to process



We can purely employ message queue


How to identify chains?

- Employ topic mode of RabbitMQ
- By marking instances with different topics






How to determine instance scale?

- Allocate more instances to urgent chain
 - Based on request arrival rate
 - Based on service times of microservices



Modeling microservice systems

- Matrix $R = (r_{c,m})_{C \times M}$
 - C is the total number of microservice chains
 - M is the number of microservice types
 - $r_{c,m} \in \{0, 1\}$ is whether chain **c** traverses microservice **m**




1	0	1	0
1	1	0	1
0	1	1	1

3×4

- ✓ Chain Red: **A -> C**
- ✓ Chain Blue: **A -> C -> D**
- ✓ Chain Green: **B -> C -> D**

Allocating instances for chains

- Matrix $S = (s_{c,m})_{C \times M}$
 - $s_{c,m}$ denotes the number of microservice m instances assigned to chain c



Evaluating response time

- Model single microservice as an M/G/1 queue
 - Request arrival follows Poisson process
 - Service time is generally distributed
 - Response time in of a microservice

$$u_{c,m}(s_{c,m}) = \frac{E[Z^{c,m}]}{1 - \rho_{c,rn}(s_{c,m})}$$


- Based on Theorem 1, single microservice is extended to a chain
 - Response time of a chain

$$u_c(s_c) = \sum_{m=1}^M r_{c,m} \cdot u_{c,n}(s_{c,m})$$

where $r_{c,m}$ indicates whether chain c passes microservice m .

Determining instance allocation

- Imagine every chain is a person
- Negotiate with each other



Game theory based instance allocation

- Utility function
 - We introduce a predefined worst response time U_c^b
 - $U_c^b = u_c(s_c^b)$ and s_c^b is the **disagreement** instance allocation of chain **c**
 - Utility function is defined as follows

$$\frac{U_c^b - u_c}{U_c^b}$$

- Utilizing Nash bargaining solution for instance allocation

$$\max_{u^c} \prod_{u^c \in U} \frac{U_c^b - u_c}{U_c^b}$$

- There exists a Nash Bargaining solution solve the above problem

COLBA: Chain-Oriented Load Balancing Algorithm

- By relaxing integrality constraints
- The original problem is reduced to convex optimization

$$\max_{\mathbf{s}_c} \quad \sum_{c=1}^C \ln \frac{U_c^b - u_c}{U_c^b}$$

Equivalent to the original problem

$$\text{s.t.} \quad \sum_{c=1}^C r_{c,m} \cdot s_{c,m} \leq I^m,$$

Instance capacity limit

$$s_{c,m} > \frac{\lambda^{c,m}}{\mu_{c,m}},$$

Ensure traffic intensity of queues

- KKT-conditions with rounding


Optimality analysis

- The optimal value of objective function: ϕ^*
- The optimal value of objective function under relaxation: $\tilde{\phi} > \phi^*$
 - The solution is $\tilde{s}_{c,m}$
- The value of objective function under relaxation with rounding ϕ
 - The solution is $\hat{s}_{c,m} = \lfloor \tilde{s}_{c,m} \rfloor$
 - So we have:
$$\phi^* - \phi < \tilde{\phi} - \phi$$
 - Finally we prove that


$$\phi^* - \phi < \sum_{c=1}^C \ln(1 + M \cdot Q_c)$$

Simulation setup

- Real world trace
 - ❑ Online traffic in U.S. on Cyber Monday measured by Akamai
- Baselines
 - ❑ Instance-oriented load balancing
 - Apply the chain-oriented load balancing strategy for fairness.
 - ❑ Microservice-oriented load balancing
- Worst response time U_c^b
 - ❑ The upper bound of response times
 - ❑ The performance requirement of user requests in chain c
 - ❑ The initial instance assignment for each chain




Request arrival rate



Worst response time $U=[10; 8; 6; 7; 2; 4; 2]$, Request rate: 1.0λ


Request arrival rate



Worst response time $U=[10; 8; 6; 7; \textcolor{red}{2}; 4; 2]$, Request rate: 1.0λ

- For Chain 5, the instances allocated to it increase significantly so as to overcome the bursty workload


Request arrival rate



Worst response time $U=[10; 8; 6; 7; \textcolor{red}{2}; 4; 2]$, Request rate: 1.5λ


- For Chain 5, the instances allocated to it increase significantly so as to overcome the bursty workload

Request arrival rate



Worst response time $U=[\mathbf{10}; 8; 6; 7; 2; 4; 2]$, Request rate: 1.0λ


Request arrival rate



Worst response time $U=[\mathbf{10}; 8; 6; 7; 2; 4; 2]$, Request rate: 1.0λ

- Number of instances assigned to Chains 1 decreases sharply, leading to increase in response time.


Request arrival rate



Worst response time $U=[\mathbf{10}; 8; 6; 7; 2; 4; 2]$, Request rate: 1.5λ


- Number of instances assigned to Chains 1 decreases sharply, leading to increase in response time.

Worst response time



Worst response time $U=[3; 3; 2; \textcolor{red}{5}; 4; 2; \textcolor{red}{9}]$, Request rate: λ


Worst response time



Worst response time $U=[3; 3; 2; \textcolor{red}{5}; 4; 2; \textcolor{red}{9}]$, Request rate: λ

- The scale of instances allocated to Chain 7 increases remarkably, especially for Microservice 1 and 6


Worst response time



Worst response time $U=[3; 3; 2; \mathbf{9}; 4; 2; \mathbf{5}]$, Request rate: λ


- The scale of instances allocated to Chain 7 increases remarkably, especially for Microservice 1 and 6

Worst response time



Worst response time $U=[3; 3; 2; \mathbf{5}; 4; 2; \mathbf{9}]$, Request rate: λ


Worst response time



Worst response time $U=[3; 3; 2; \textcolor{red}{5}; 4; 2; \textcolor{red}{9}]$, Request rate: λ

- Instance scale of Microservice 6 assigned to Chain 4 shrinks significantly

Worst response time



Worst response time $U=[3; 3; 2; \textcolor{red}{9}; 4; 2; \textcolor{red}{5}]$, Request rate: λ

- Instance scale of Microservice 6 assigned to Chain 4 shrinks significantly

Conclusion

- We proposed COLBA to balance load from perspective of chains
- Leveraging message queues, requests can be efficiently balanced across microservice instances
- Simulation results show that COLBA can coordinate instance allocation, concerning expected QoS of chains as well as competition across chains

Q&A

Thank You!

“Cloud Datacenter & Green Computing” Research Group
Huazhong University of Science & Technology

<http://grid.hust.edu.cn/fmliu/>
fmliu@hust.edu.cn

Discussion & open problems

- Q: What are the differences between microservice chain and NFV chain?
 - A: Chain path is implemented in different layers.
- Q: Is it OK to abandon HTTP?
 - A: Modern web applications are developed using the Task Asynchronous Paradigm*.
 - The states of connections can also be enabled with message queue.

*Webperf: Evaluating what-if scenarios for cloud-hosted web applications," in Proc. of SIGCOMM, 2016.